

Approximating Sequence Similarity

Brad Solomon

March 27, 2019

Lecture 16: Scalable Methods for Genomics



Approximating Sequence Similarity

Brad Solomon

With Special Focus on RNA-Seq

March 27, 2019

Lecture 16: Scalable Methods for Genomics



A Quick Recap

Overlap between two sequences

...AGCCTAGACCTACAGGATGCGCGGACACGTAGCCAGGAC
CAGTACTTGGATGCGCTGACACGTAGCTTATCCGGT...

overlap (19 bases) overhang (6 bases)

overhang % identity = 18/19 % = 94.7%

overlap - region of similarity between regions

overhang - un-aligned ends of the sequences

The assembler screens merges based on:


- length of overlap
- % identity in overlap region
- maximum overhang size.

[How do we compute the overlap?]

[Do we really want to do all-vs-all?]

See Lecture 4 Assembly & Whole
Genome Alignment

RNA-seq Challenges




Challenge 1: Eukaryotic genes are spliced

Solution: Use a spliced aligner, and assemble isoforms

TopHat: discovering spliced junctions with RNA-Seq.

Trapnell et al (2009) *Bioinformatics*. 25:0 1105-1111




Challenge 2: Read Count != Transcript abundance

Solution: Infer underlying abundances (e.g. TPM)

Transcript assembly and quantification by RNA-seq

Trapnell et al (2010) *Nat. Biotech.* 25(5): 511-515



Challenge 3: Transcript abundances are stochastic


Solution: Replicates, replicates, and more replicates

RNA-seq differential expression studies: more sequence or more replication?


Liu et al (2013) *Bioinformatics*. doi:10.1093/bioinformatics/btt688

See Lecture 10 RNA Sequencing

Sequence similarity at scale is a ‘big’ problem



Sequence similarity at scale is a ‘big’ problem



Biological questions at scale

Biological questions at scale

Experiment Discovery

I have sample X, find me related samples

Biological questions at scale

Experiment Discovery

I have sample X, find me related samples

Containment Query

I am interested in transcript X, find me studies expressing it

Biological questions at scale

Experiment Discovery

I have sample X, find me related samples

Containment Query

I am interested in transcript X, find me studies expressing it

Cardinality Estimation

I have N studies, how many unique sequences are present?

Biological questions at scale

Experiment Discovery

I have sample X, find me related samples

Containment Query

I am interested in transcript X, find me studies expressing it

Cardinality Estimation

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns
of transcriptome X on N studies

Biological questions at scale

Experiment Discovery

I have sample X, find me related samples

Containment Query

I am interested in transcript X, find me studies expressing it

Cardinality Estimation

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns of transcriptome X on N studies

How can we solve these problems?

Biological questions at scale

Experiment Discovery

I have sample X, find me related samples

Containment Query

I am interested in transcript X, find me studies expressing it

Cardinality Estimation

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns of transcriptome X on N studies

How can we solve these problems?

Biological questions at scale

Experiment Discovery Measure similarity as **alignment** overlap

I have sample X, find me related samples

Containment Query

I am interested in transcript X, find me studies expressing it

Cardinality Estimation

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns
of transcriptome X on N studies

How can we solve these problems?

Biological questions at scale

Experiment Discovery Measure similarity as **alignment** overlap

I have sample X, find me related samples

Containment Query **Align** each study to a reference, observe X

I am interested in transcript X, find me studies expressing it

Cardinality Estimation

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns
of transcriptome X on N studies

How can we solve these problems?

Biological questions at scale

Experiment Discovery Measure similarity as **alignment** overlap

I have sample X, find me related samples

Containment Query **Align** each study to a reference, observe X

I am interested in transcript X, find me studies expressing it

Align each study to a reference, observe **Cardinality Estimation**

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns
of transcriptome X on N studies

How can we solve these problems?

Biological questions at scale

Experiment Discovery Measure similarity as **alignment** overlap

I have sample X, find me related samples

Containment Query **Align** each study to a reference, observe X

I am interested in transcript X, find me studies expressing it

Align each study to a reference, observe **Cardinality Estimation**

I have N studies, how many unique sequences are present?

Align each study to a reference, normalize and observe **Expression Estimation**

I am interested in the global expression patterns
of transcriptome X on N studies

How can we solve these problems?

Methods for alignment at scale

ExAC Browser Beta

 **Rail-RNA**

 **recount2**

Methods for alignment at scale

ExAC Browser Beta

 Rail-RNA

 recount2



Methods for alignment at scale

ExAC Browser Beta

 Rail-RNA

 recount2



Rail-RNA: Bulk Alignment

Given a set of N sequencing studies,

- 1) **Aggregate** the reads into sets of overlapping **reads** and **readlets** (a subsequence of a set of reads with partial overlap)
- 2) Perform **parallel alignment** where each thread is given distinct **reads** based on nucleotide similarity

Rail-RNA: Bulk Alignment

Given a set of N sequencing studies,

- 1) **Aggregate** the reads into sets of overlapping **reads** and **readlets** (a subsequence of a set of reads with partial overlap)
- 2) Perform **parallel alignment** where each thread is given distinct **reads** based on nucleotide similarity
- 3) **Second pass alignment** on all non-perfect or tied score alignments adds quality scores to break ties. Additional **readlets** are generated for poor alignments
- 4) **Readlets** are **parallel aligned** based on nucleotide similarity


Rail-RNA: Bulk Alignment

Given a set of N sequencing studies,

- 1) **Aggregate** the reads into sets of overlapping **reads** and **readlets** (a subsequence of a set of reads with partial overlap)
- 2) Perform **parallel alignment** where each thread is given distinct **reads** based on nucleotide similarity
- 3) **Second pass alignment** on all non-perfect or tied score alignments adds quality scores to break ties. Additional **readlets** are generated for poor alignments
- 4) **Readlets** are **parallel aligned** based on nucleotide similarity
- 5) Further steps are taken to identify **exon-exon** junctions and one final alignment is performed using the **bulk exon data** identified from the studies
- 6) All individual alignments are output, with a single primary alignment compiled for each individual study

Recount2: Combining methods

1) A collection of studies is selected



48,558 samples from SRA
11,350 samples from TCGA
9,662 samples from GTEx

2) Rail-RNA was run on cloud computing services in batches



BigWig


3) Each alignment is stored as a BigWig
(a dense data storage structure)

4) The resulting data was made publicly available



Recount2: Combining methods

1) A collection of studies is selected



48,558 samples from SRA
11,350 samples from TCGA
9,662 samples from GTEx

How were these selected?

2) Rail-RNA was run on cloud computing services in batches



BigWig

3) Each alignment is stored as a BigWig
(a dense data storage structure)


4) The resulting data was made publicly available



Recount2: Combining methods

How were the batches selected?

1) A collection of studies is selected



2) Rail-RNA was run on cloud computing services in batches



48,558 samples from SRA
11,350 samples from TCGA
9,662 samples from GTEx

How were these selected?

BigWig

4) The resulting data was made publicly available




3) Each alignment is stored as a BigWig
(a dense data storage structure)

Recount2: Combining methods

How were the batches selected?

1) A collection of studies is selected



2) Rail-RNA was run on cloud computing services in batches




48,558 samples from SRA
11,350 samples from TCGA
9,662 samples from GTEx

How is this designed?

4) The resulting data was made publicly available


How were these selected?

BigWig




3) Each alignment is stored as a BigWig
(a dense data storage structure)


Alignment can't scale to Big Data



Alignment can't scale to Big Data



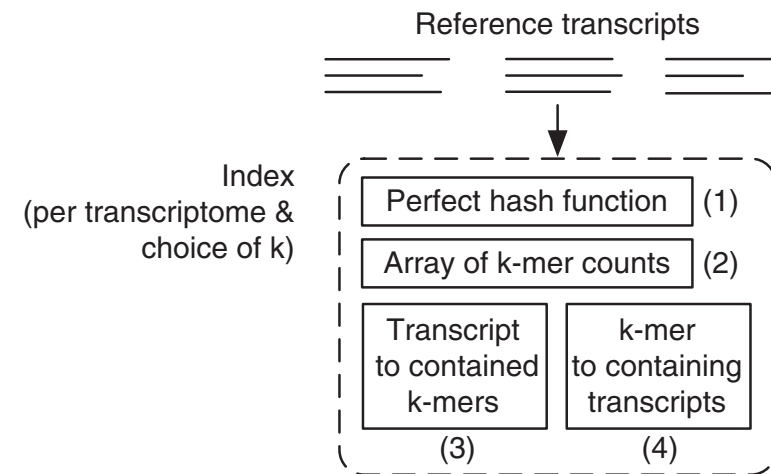
Alignment can't scale to Big Data



Recap: Sailfish

1) Indexing

- Parse reference transcriptome into kmers
- All unique kmers are hashed and counted
- Two indices for bi-directional mapping of transcripts and kmers



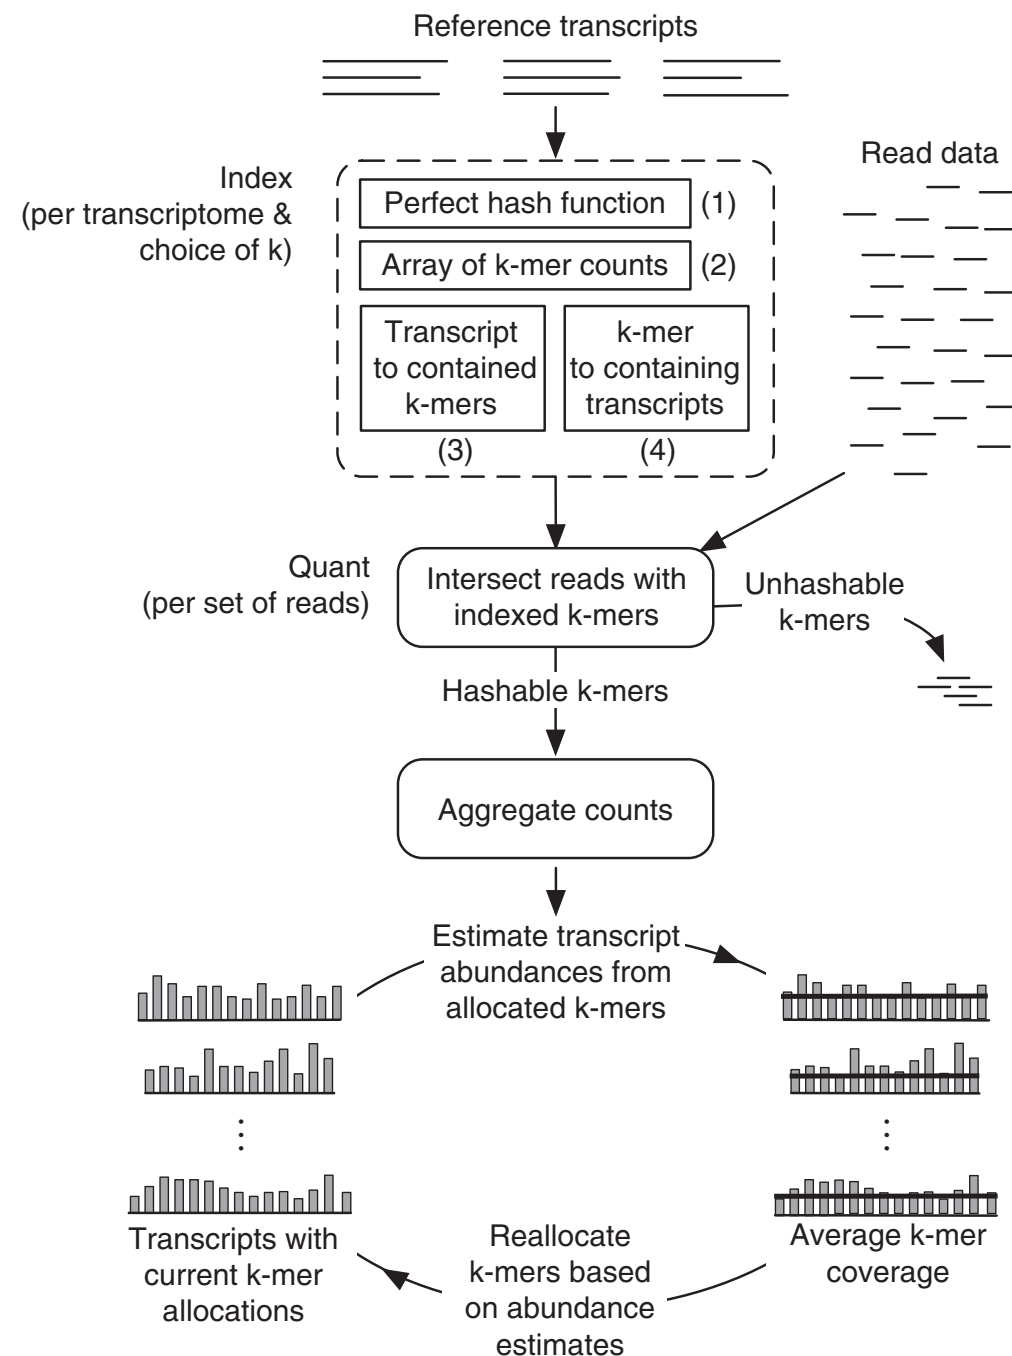
Recap: Sailfish

1) Indexing

- Parse reference transcriptome into kmers
- All unique kmers are hashed and counted
- Two indices for bi-directional mapping of transcripts and kmers

2) Quantification

- Count kmers in read set
- Use EM procedure to estimate transcript abundances, repeating as necessary



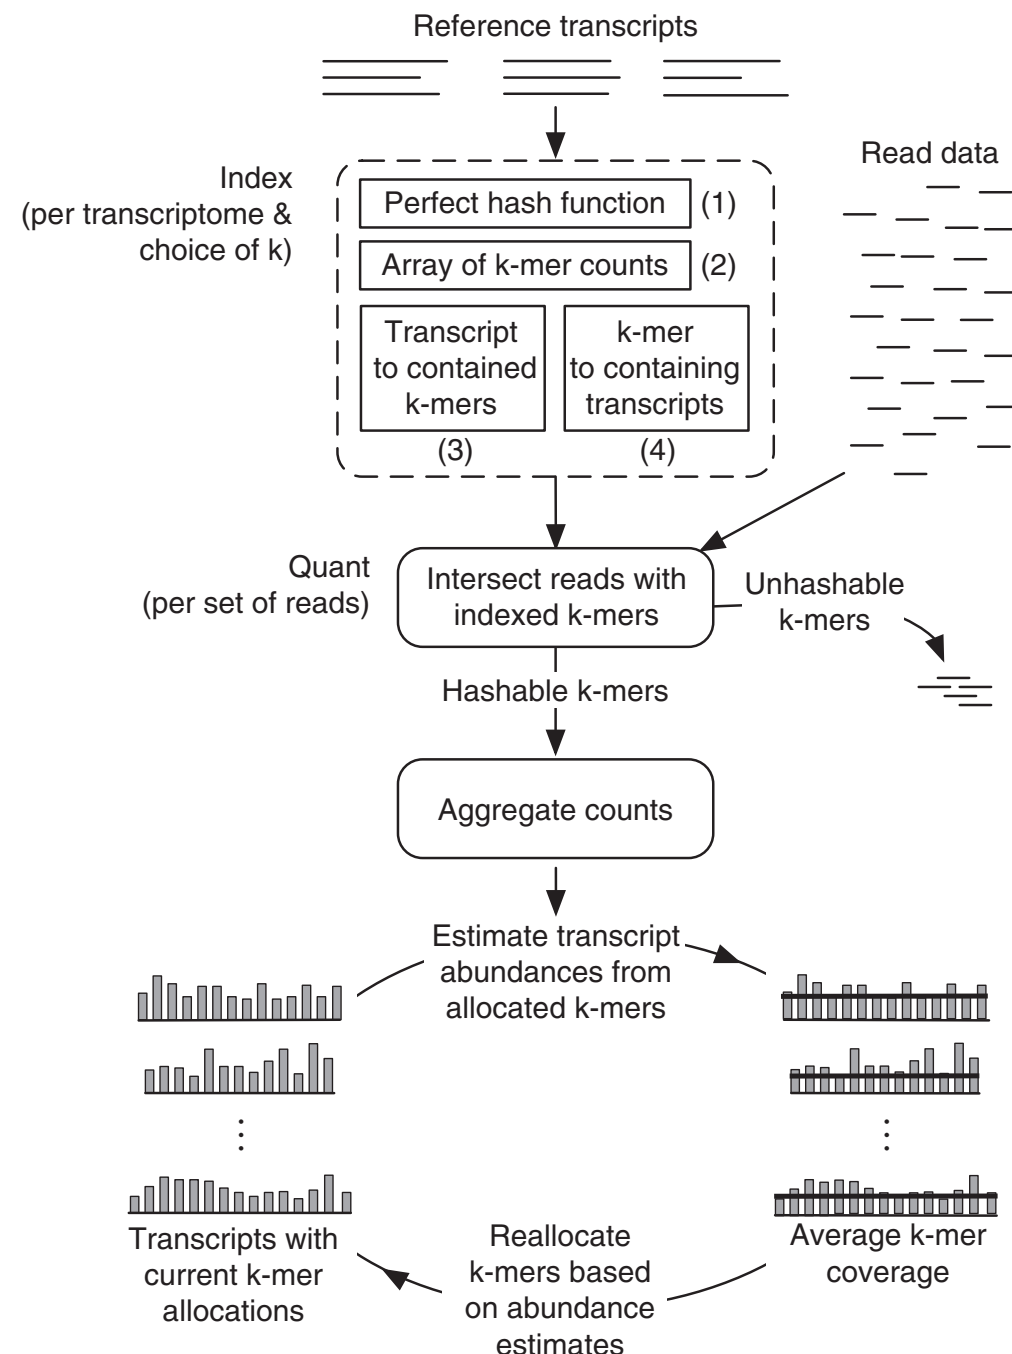
Recap: Sailfish

1) Indexing

- Parse reference transcriptome into kmers
- All unique kmers are hashed and counted
- Two indices for bi-directional mapping of transcripts and kmers

2) Quantification

- Count kmers in read set
- Use EM procedure to estimate transcript abundances, repeating as necessary



The Big Data Problem

- There's an abundance of underutilized sequence data
- Alignment can solve a number of important biological questions but doesn't scale
- Alignment-free methods are more efficient but are not *several orders of magnitude* more efficient

The Big Data Problem

- There's an abundance of underutilized sequence data
- Alignment can solve a number of important biological questions but doesn't scale
- Alignment-free methods are more efficient but are not *several orders of magnitude* more efficient

How can we address big data?

The Big Data Problem

- There's an abundance of underutilized sequence data
- Alignment can solve a number of important biological questions but doesn't scale
- Alignment-free methods are more efficient but are not *several orders of magnitude* more efficient

How can we address big data?

Sketching!

Sketching algorithms trade accuracy for speed

Given a box, it's easy to
discover what's inside



Sketching algorithms trade accuracy for speed

Given a box, it's easy to discover what's inside



But what if there's too many boxes?



Sketching algorithms trade accuracy for speed

Given a box, it's easy to discover what's inside



But what if there's too many boxes?



A sketch solution would organize boxes based on labels, simpler observations such as size or weight, or sub-sample boxes to learn about its neighbors

Tradeoffs in Similarity Metrics

- Hamming distance
 - Count the number of substitutions to transform one string into another

MIKESCHATZ

| | x | | xxxx |

MICESHATZZ

5

- Edit distance
 - The minimum number of substitutions, insertions, or deletions to transform one string into another

MIKESCHAT-Z

| | x | | x | | | x |

MICES-HATZZ

3

Tradeoffs in Similarity Metrics

- Hamming distance
 - Count the number of substitutions to transform one string into another

MIKESCHATZ

| | x | | xxxx |

MICESHATZZ

5

Much faster to calculate

- Edit distance
 - The minimum number of substitutions, insertions, or deletions to transform one string into another

MIKESCHAT-Z

| | x | | x | | | x |

MICES-HATZZ

3

Tradeoffs in Similarity Metrics

- Hamming distance
 - Count the number of substitutions to transform one string into another

MIKESCHATZ

| | X | | XXXX |

MICESHATZZ

5

Much faster to calculate

- Edit distance
 - The minimum number of substitutions, insertions, or deletions to transform one string into another

MIKESCHAT-Z

| | X | | X | | | X |

MICES-HATZZ

3

More biologically meaningful

Tradeoffs in Similarity Metrics

- Hamming distance
 - Count the number of substitutions to transform one string into another

MIKESCHATZ

| | X | | XXXX |

MICESHATZZ

5

Much faster to calculate

- Edit distance
 - The minimum number of substitutions, insertions, or deletions to transform one string into another

MIKESCHAT-Z

| | X | | X | | | X |

MICES-HATZZ

3

More biologically meaningful

Can you think of other tradeoffs in alignment methods?

Solving Big Data through Sketching

Experiment Discovery

I have sample X, find me related samples

Containment Query

I am interested in transcript X, find me experiments expressing it

Cardinality Estimation

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns of transcriptome X in N studies

Solving Big Data through Sketching

Experiment Discovery Mash [Minhash]

I have sample X, find me related samples

Containment Query

I am interested in transcript X, find me experiments expressing it

Cardinality Estimation

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns of transcriptome X in N studies

Solving Big Data through Sketching

Experiment Discovery Mash [Minhash]

I have sample X, find me related samples

Containment Query Sequence Bloom Trees [Bloom Filter]

I am interested in transcript X, find me experiments expressing it

Cardinality Estimation

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns of transcriptome X in N studies

Solving Big Data through Sketching

Experiment Discovery Mash [Minhash]

I have sample X, find me related samples

Containment Query Sequence Bloom Trees [Bloom Filter]

I am interested in transcript X, find me experiments expressing it

Dashing [HyperLogLog] **Cardinality Estimation**

I have N studies, how many unique sequences are present?

Expression Estimation

I am interested in the global expression patterns
of transcriptome X in N studies

Solving Big Data through Sketching

Experiment Discovery Mash [Minhash]

I have sample X, find me related samples

Containment Query Sequence Bloom Trees [Bloom Filter]

I am interested in transcript X, find me experiments expressing it

Dashing [HyperLogLog] **Cardinality Estimation**

I have N studies, how many unique sequences are present?

Salmon / Kallisto **Expression Estimation**

I am interested in the global expression patterns
of transcriptome X in N studies

Solving Big Data through Sketching

Experiment Discovery *Mash [Minhash]*

I have sample X, find me related samples

Containment Query *Sequence Bloom Trees [Bloom Filter]*

I am interested in transcript X, find me experiments expressing it

Dashing [HyperLogLog] **Cardinality Estimation**

I have N studies, how many unique sequences are present?

Salmon / Kallisto **Expression Estimation**

I am interested in the global expression patterns
of transcriptome X in N studies

Mantis [Counting Quotient Filter]

The Minhash Sketch

1) Sequence decomposed
into **kmers**

S_1 : CATGGACCGACCAG
CAT GAC GAC
ATG ACC ACC
TGG CCG CCA
GGA CGA CAG

GCAGTACCGATCGT : S_2
GTA CGA CGT
AGT CCG TCG
CAG ACC ATC
GCA TAC GAT

The Minhash Sketch

1) Sequence decomposed into **kmers**

2) Multiple hash functions (Γ) map kmers to values.

S_1 : CATGGACCGACCAG
 CAT GAC GAC
 ATG ACC ACC
 TGG CCG CCA
 GGA CGA CAG

Γ_1	Γ_2	Γ_3	Γ_4	
19	14	57	36	CAT
14	57	36	19	ATG
58	37	16	15	TGG
40	23	2	61	GGA
33	28	11	54	GAC
5	48	47	26	ACC
22	1	60	43	CCG
24	7	50	45	CGA
33	28	11	54	GAC
5	48	47	26	ACC
20	3	62	41	CCA
18	13	56	39	CAG

GCAGTACCGATCGT : S_2
 GTA CGA CGT
 AGT CCG TCG
 CAG ACC ATC
 GCA TAC GAT

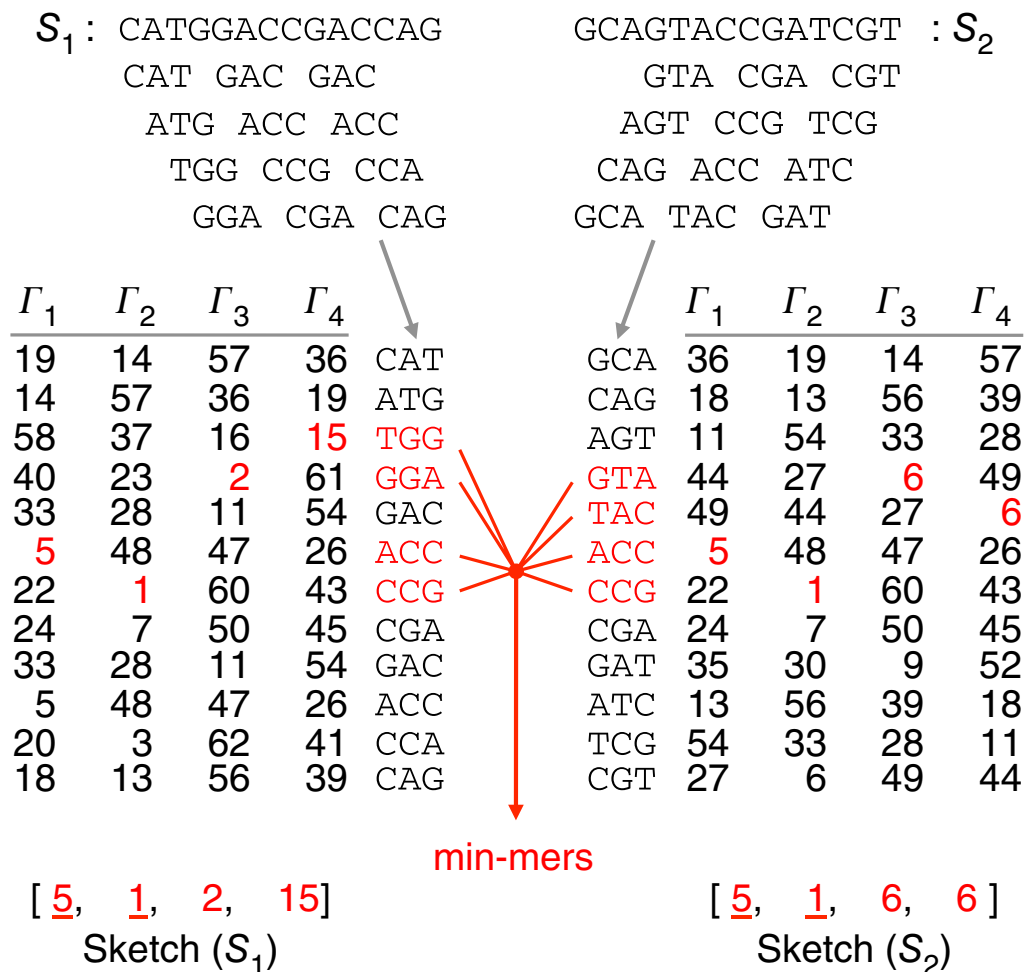
	Γ_1	Γ_2	Γ_3	Γ_4
GCA	36	19	14	57
CAG	18	13	56	39
AGT	11	54	33	28
GTA	44	27	6	49
TAC	49	44	27	6
ACC	5	48	47	26
CCG	22	1	60	43
CGA	24	7	50	45
GAT	35	30	9	52
ATC	13	56	39	18
TCG	54	33	28	11
CGT	27	6	49	44

The Minhash Sketch

1) Sequence decomposed into **kmers**

2) Multiple hash functions (Γ) map kmers to values.

3) The smallest values for each hash function is chosen



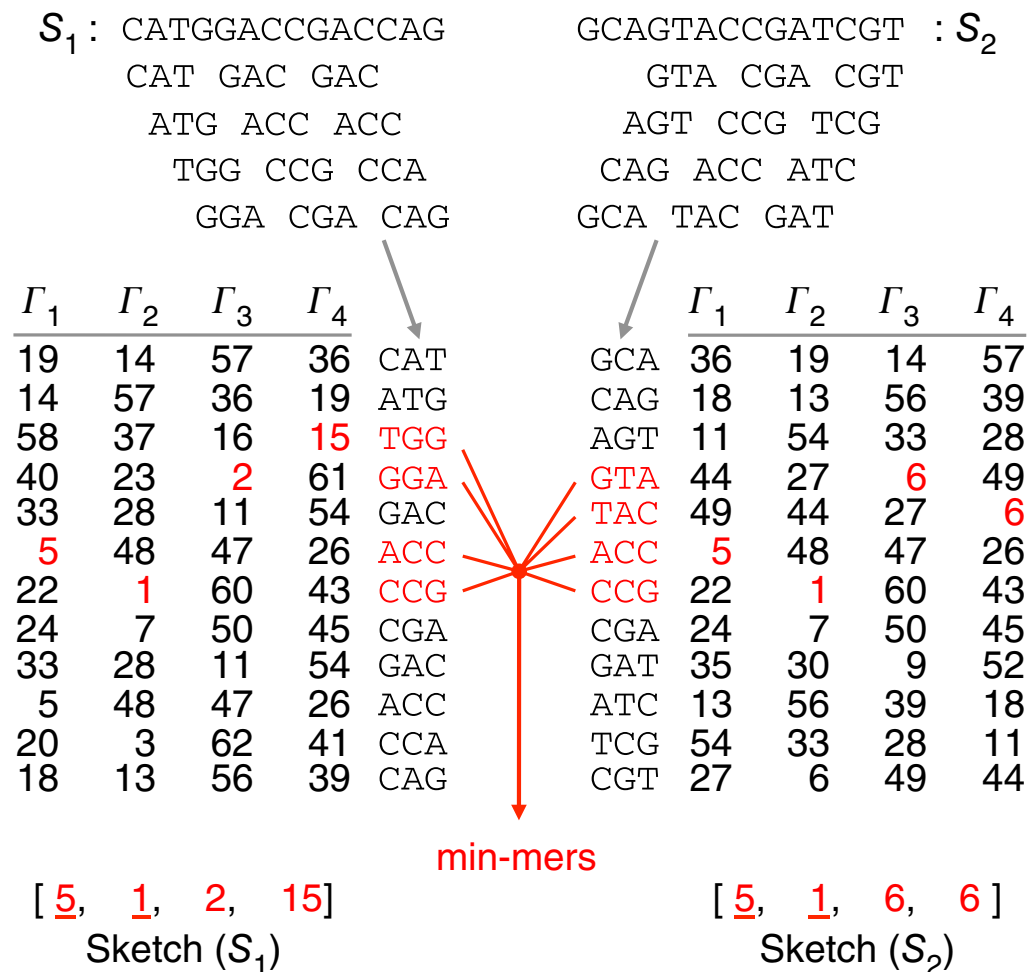
The Minhash Sketch

1) Sequence decomposed into **kmers**

2) Multiple hash functions (Γ) map kmers to values.

3) The smallest values for each hash function is chosen


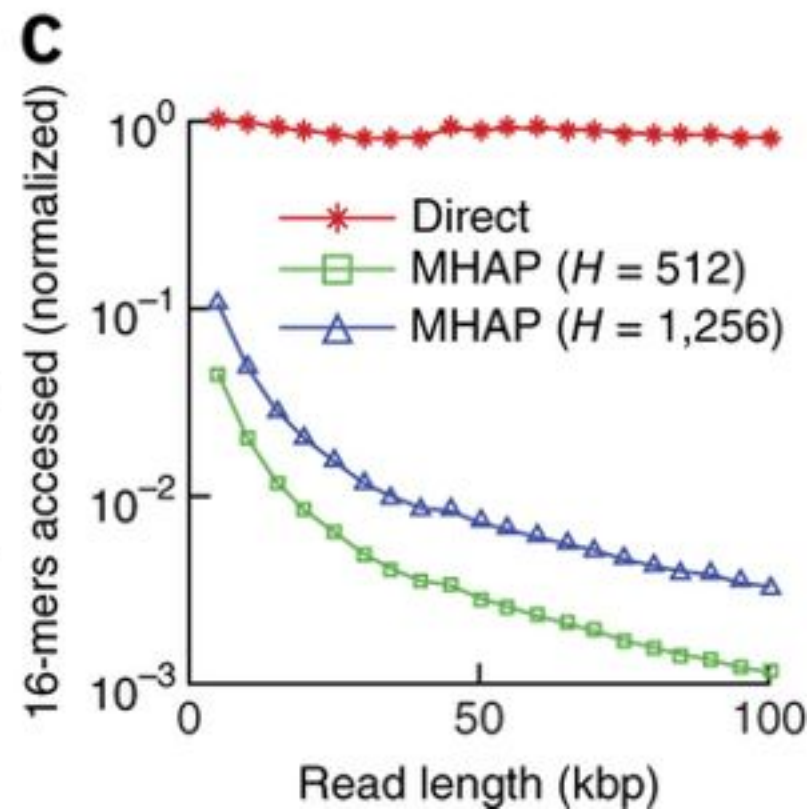
4) The Jaccard similarity can be estimated by the overlap in the **Minimum Hashes (Minhash)**



$$J(S_1, S_2) \approx 2/4 = 0.5$$


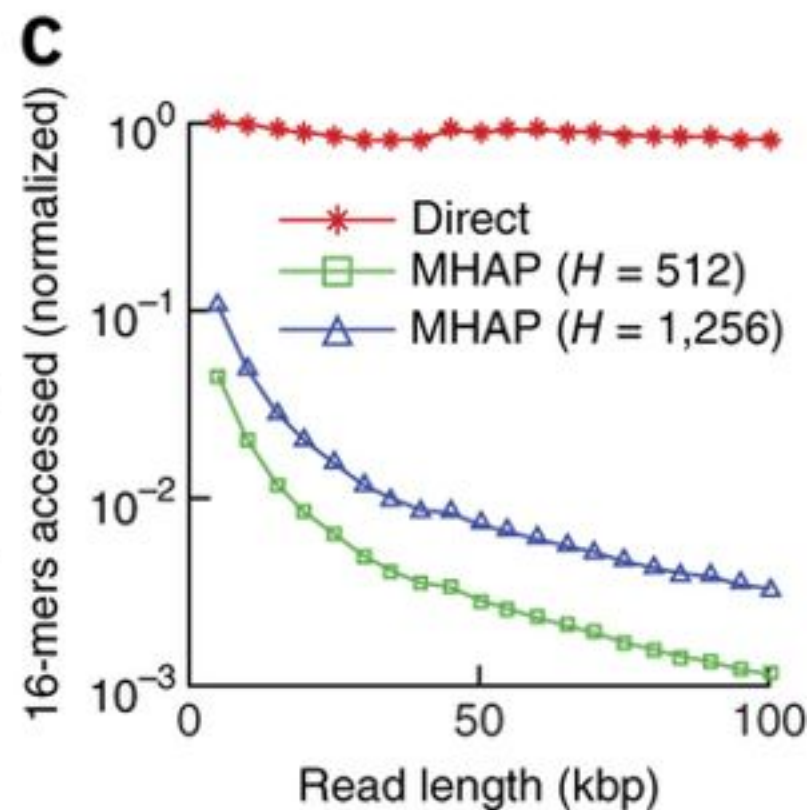
S_1 : CATGGACCGACCAG
 | | | | |
 S_2 : GCAGTACCGATCGT

MHAP uses Minhash to approximate read overlap

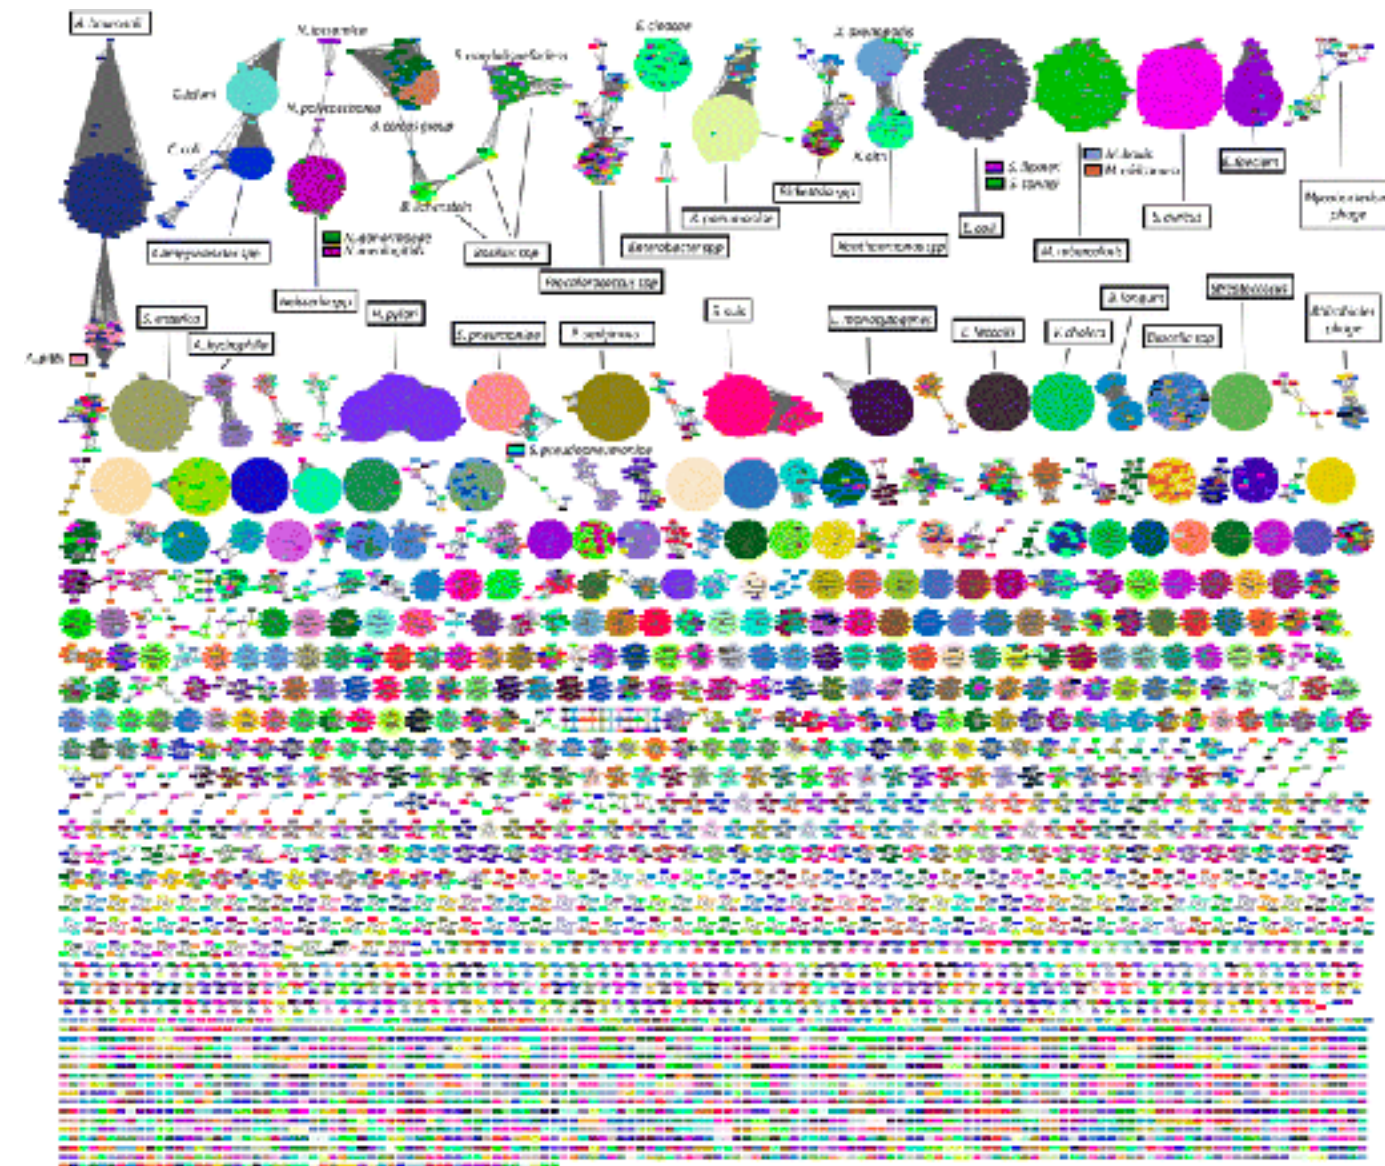


MHAP uses Minhash to approximate read overlap

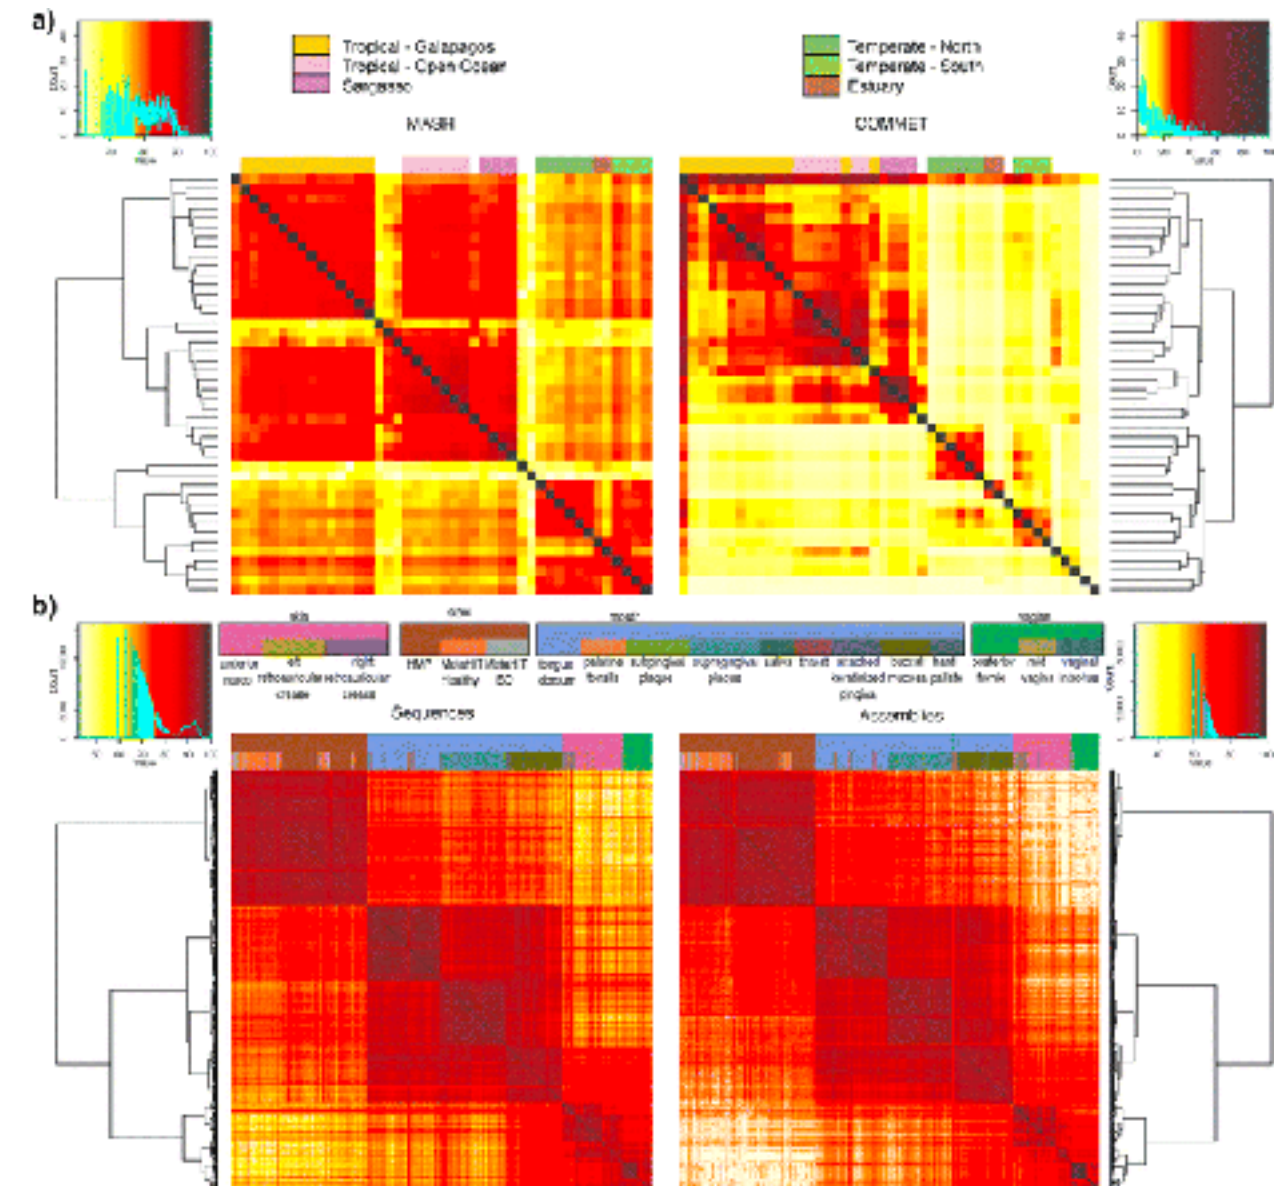
An improvement in heuristic efficiency leads to an improvement in accuracy



Mash uses Minhash to approximate genome and read overlap




Ref Seq Genome Clusterings



Read set overlap

Bloom Filters efficiently encode sets

A Bloom Filter is a length m bit-vector and associated hash function(s) $H(X)$




Bloom Filters efficiently encode sets

A Bloom Filter is a length m bit-vector and associated hash function(s) $H(X)$

Hash Function $H(X)$ takes in an arbitrary element and returns an integer

Ex: $H(\text{[redacted]}) = 12$




Bloom Filters efficiently encode sets

A Bloom Filter is a length m bit-vector and associated hash function(s) $H(X)$

Hash Function $H(X)$ takes in an arbitrary element and returns an integer

Ex: $H(\text{blue box}) = 12$




Bloom Filters efficiently encode sets

A Bloom Filter is a length m bit-vector and associated hash function(s) $H(X)$

Hash Function $H(X)$ takes in an arbitrary element and returns an integer

Ex: $H(\text{blue box}) = 12$

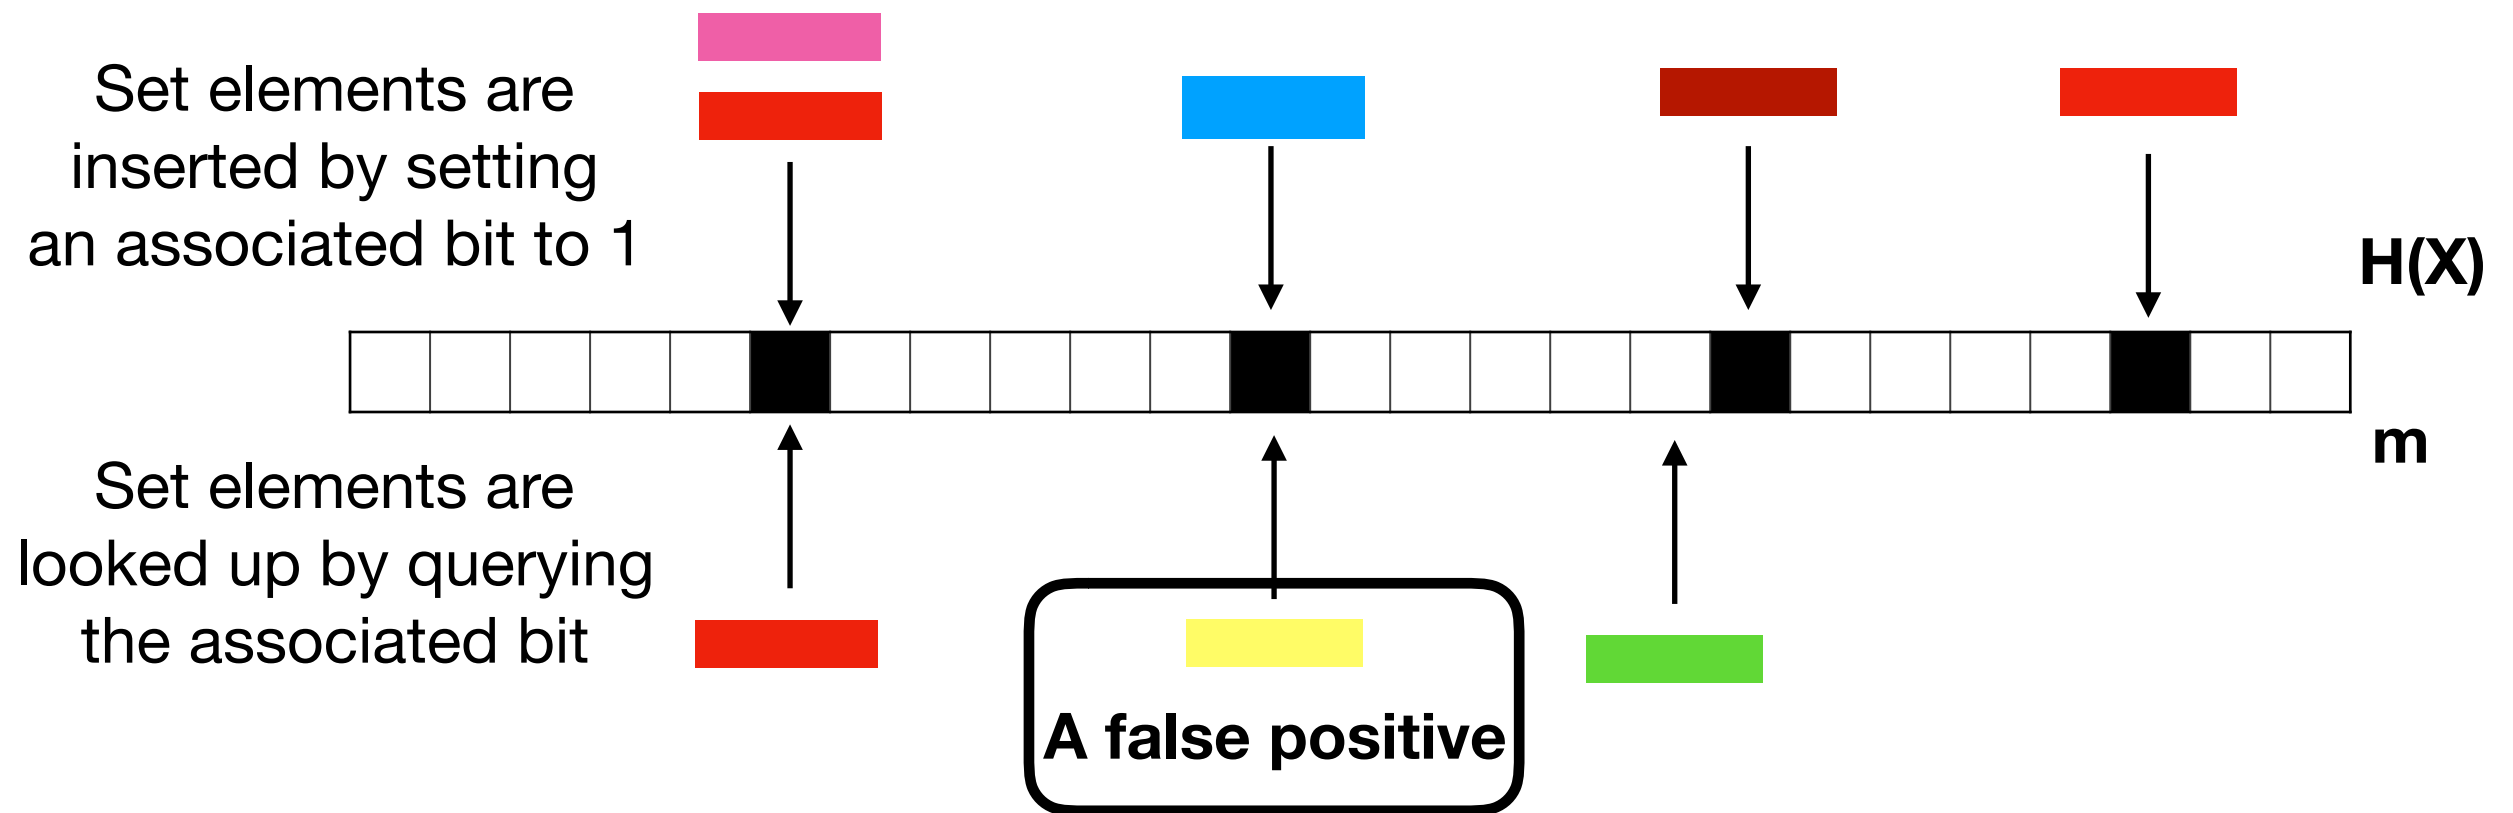


Bloom Filters efficiently encode sets

A Bloom Filter is a length m bit-vector and associated hash function(s) $H(X)$

Hash Function $H(X)$ takes in an arbitrary element and returns an integer

Ex: $H(\text{blue}) = 12$




A Bloom Filter is probabilistic data structure!

Bloom Filters efficiently encode sets

A Bloom Filter is a length m bit-vector and associated hash function(s) $H(X)$

Hash Function $H(X)$ takes in an arbitrary element and returns an integer

Ex: $H(\text{blue}) = 12$




A Bloom Filter is probabilistic data structure!

Stores arbitrary sets and supports $O(1)$ insertion and membership testing

The “Containment” Query

Input:

- Set of individual sequencing studies




Each study contains a set of raw reads

The “Containment” Query

Input:

- Set of individual sequencing studies



- A query of interest

```
ATGGTTAGAATTAAACCTGGATC
TGCTAATAAACCUAGTGATGATG
CGATAGCACAGGTAGATCCAGT
TACGTAGAGGTCATTAGCCGTAT
TGCTAATAAACCTAGTGATGATT
CGATAGCGTAGAGGTCATTAGC
CTTGTGCTAATAACAGGTAGA
TCCGTATACGTAGAGGTCATTA
CCTTGTGCTAATAACCTAGTG
```


Ex: A novel transcript

- θ , the definition of containment

The “Containment” Query

Input:

- Set of individual sequencing studies



Ex: TCGA, NIH SRA

- A query of interest

```
ATGGTTAGAATTAAACCTGGATC
TGCTAATAAACCUAGTGATGATG
CGATAGCACAGGTAGATCCAGT
TACGTAGAGGTCATTAGCCGTAT
TGCTAATAAACCTAGTGATGATT
CGATAGCGTAGAGGTCATTAGC
CTTGTGCTAATAACAGGTAGA
TCCGTATACGTAGAGGTCATTA
CCTTGTGCTAATAACCTAGTG
```

Ex: A novel transcript

```
ATGGTTAGAATTAAACCCGG
TGCTAATAAACCUAGTGATG

CGATAGCACAGGTAGATCC
TACGTAGAGGTCATTAGCC

TACGTAGAGGTCATTAGCCG
TGCTAATAAACCUAGTGATG


....
```

Each study contains a set of raw reads

- θ , the definition of containment

Output:


All studies whose **read sets** can cover θ fraction of the query



Finding **relevant experiments** is the first step in many large-scale analyses.


**All studies whose
reads cover the query**

Output



Finding **relevant experiments** is the first step in many large-scale analyses.

All studies whose reads cover the query




Functional Enrichment
Analysis

Output

Gene enrichment
associated with a query

Finding **relevant experiments** is the first step in many large-scale analyses.

All studies whose reads cover the query



Functional Enrichment
Analysis

Output


Gene enrichment
associated with a query

Variant Calling

Identify novel SNPs
within a population

Finding **relevant experiments** is the first step in many large-scale analyses.

All studies whose reads cover the query



Functional Enrichment
Analysis

Variant Calling

Expression Estimation


Output

Gene enrichment
associated with a query

Identify novel SNPs
within a population

Expression levels in a
population

The Sequence Bloom Tree



The Sequence Bloom Tree

1) Sequence reads are broken down into kmers


ATGGTTAGAATTAAA

ATG TTA AAT AAA

TGG TAG ATT

GGT AGA TTA

GTT GAA TAA



The Sequence Bloom Tree

1) Sequence reads are broken down into kmers

ATGGTTAGAATTAAA



ATG TTA AAT AAA

TGG TAG ATT

GGT AGA TTA

GTT GAA TAA

2) The set of all kmers in a study are stored in a **bloom filter**





The Sequence Bloom Tree

1) Sequence reads are broken down into kmers

ATGGTTAGAATTAAA
ATG TTA AAT AAA
TGG TAG ATT
GGT AGA TTA
GTT GAA TAA

2) The set of all kmers in a study are stored in a **bloom filter**




The Sequence Bloom Tree


1) Sequence reads are broken down into kmers

ATGGTTAGAATTAAA
ATG TTA AAT AAA
TGG TAG ATT
GGT AGA TTA
GTT GAA TAA

2) The set of all kmers in a study are stored in a **bloom filter**




4) Proving that this node can't contain the query "prunes" the tree



3) Nodes store the total sequence content of all leaves rooted at that node


Conceptual Construction

Bit-wise union of bloom filter produces bloom filter with total kmer content of each file




Conceptual Construction

Bit-wise union of bloom filter produces bloom filter with total kmer content of each file



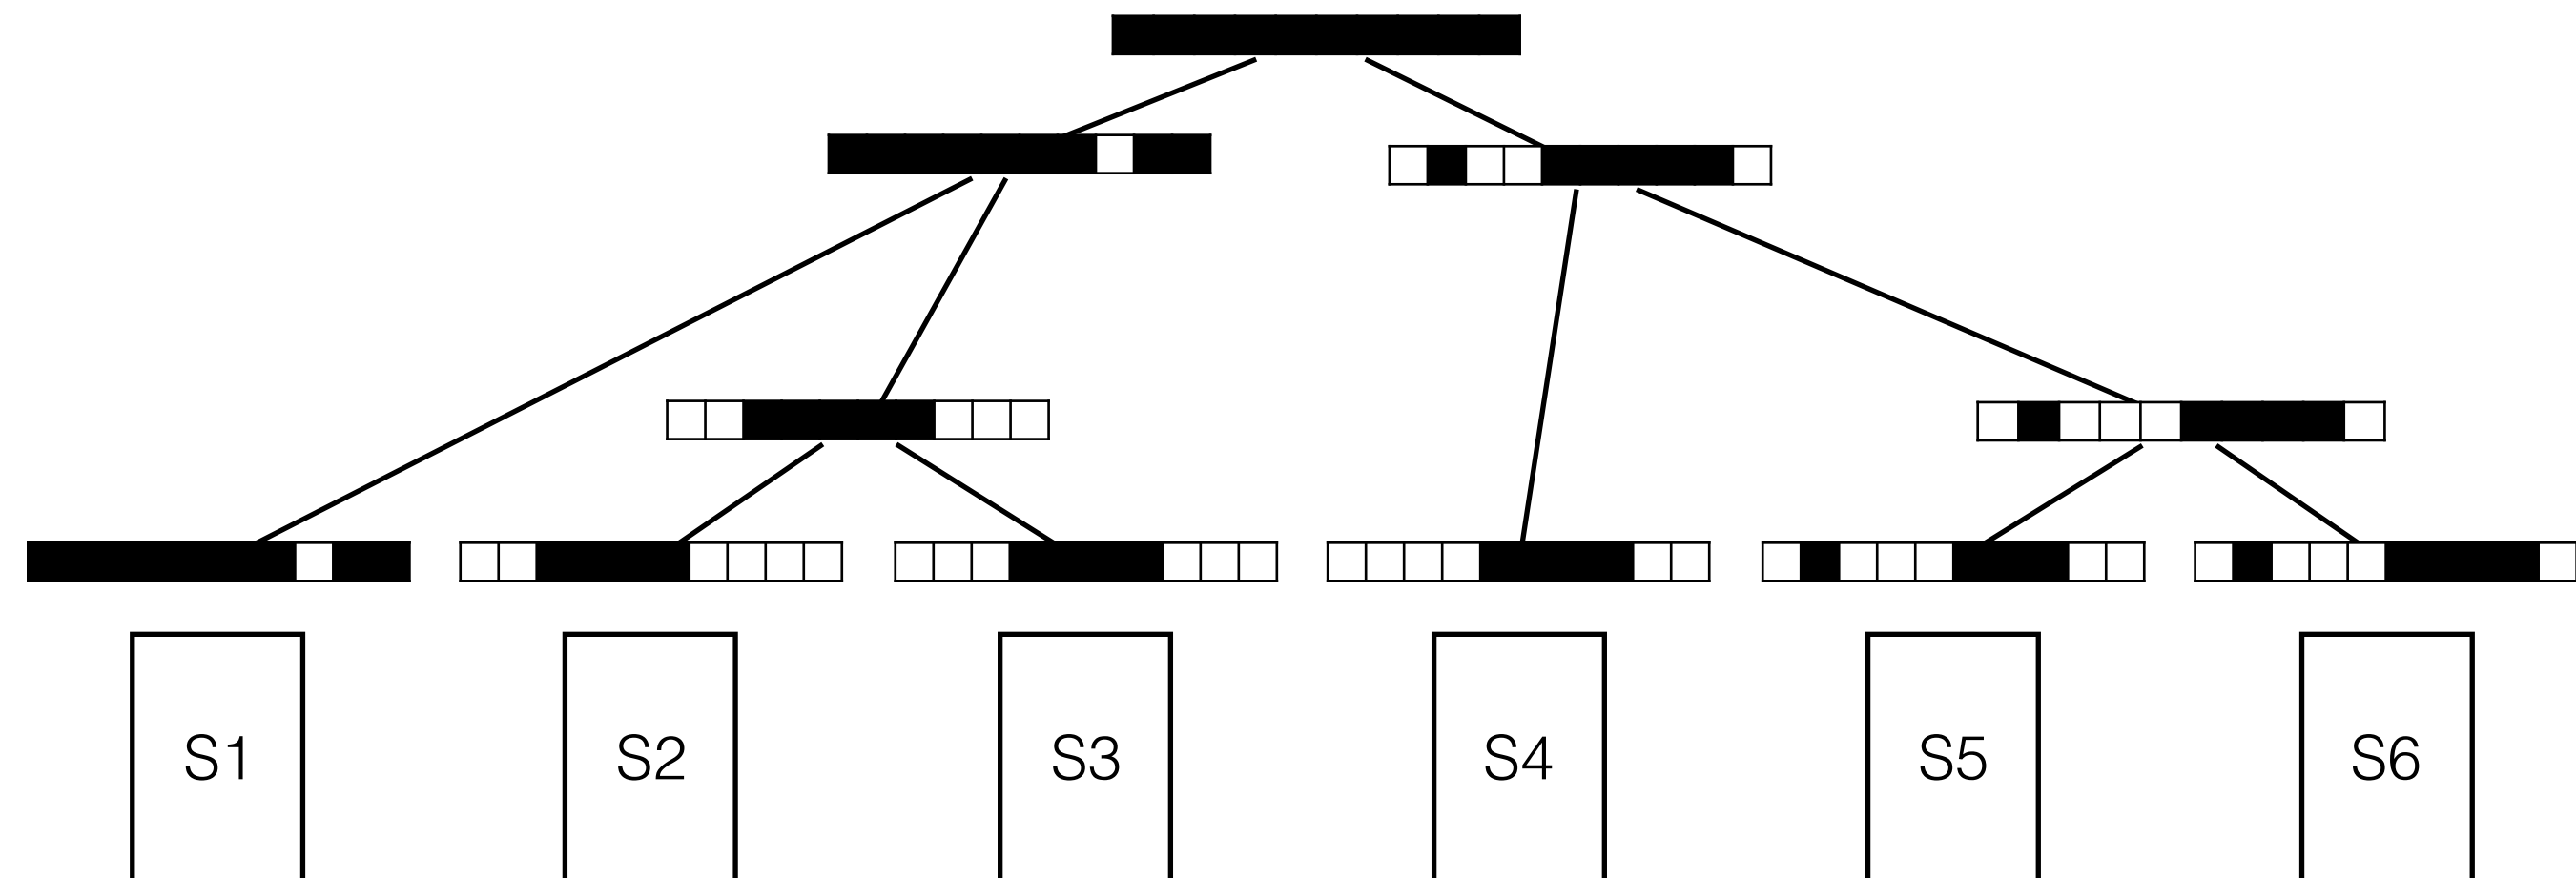
Conceptual Construction

Bit-wise union of bloom filter produces bloom filter with total kmer content of each file




Conceptual Construction

Bit-wise union of bloom filter produces bloom filter with total kmer content of each file



Similarity Metrics for Bloom Filters


Hamming distance: Number of 'substitutions' to transform one vector to another



Hamming Distance = 4



Jaccard similarity: Intersection of 1-bits over **union** of 1-bits



Intersection = 1




Union = 5



Similarity = $1/5$

Similarity Metrics for Bloom Filters

Hamming distance: Number of 'substitutions' to transform one vector to another




Hamming Distance = 4



Compare these metrics by **normalizing** (and **inverting**)

Jaccard similarity: Intersection of 1-bits over union of 1-bits



Intersection = 1




Union = 5



Similarity = $1/5$

Similarity Metrics for Bloom Filters

Hamming distance: Number of 'substitutions' to transform one vector to another




Hamming Distance = 4



Compare these metrics by **normalizing** (and **inverting**)

$$\text{Distance} / \text{Length} = 4 / 10$$

Jaccard similarity: Intersection of 1-bits over **union** of 1-bits



Intersection = 1




Union = 5



$$\text{Similarity} = 1/5$$

Similarity Metrics for Bloom Filters

Hamming distance: Number of 'substitutions' to transform one vector to another



Hamming Distance = 4




Compare these metrics by **normalizing** (and **inverting**)

$$\text{Distance} / \text{Length} = 4 / 10$$

$$(\text{Similarity} = 6/10)$$

Jaccard similarity: Intersection of 1-bits over union of 1-bits



Intersection = 1




Union = 5



$$\text{Similarity} = 1/5$$

Similarity Metrics for Bloom Filters

Hamming distance: Number of 'substitutions' to transform one vector to another



Hamming Distance = 4




Compare these metrics by **normalizing** (and **inverting**)

$$\text{Distance} / \text{Length} = 4 / 10$$

$$(\text{Similarity} = 6/10)$$

Jaccard similarity: Intersection of 1-bits over union of 1-bits



Intersection = 1



Union = 5




$$\text{Similarity} = 1/5$$

The **Hamming** similarity does not correlate with sequence content

The **Jaccard** similarity correlates with sequence content

Similarity Metrics for Bloom Filters

Hamming distance: Number of 'substitutions' to transform one vector to another



Hamming Distance = 4




Compare these metrics by **normalizing** (and **inverting**)

$$\text{Distance} / \text{Length} = 4 / 10$$

$$(\text{Similarity} = 6/10)$$

Jaccard similarity: Intersection of 1-bits over union of 1-bits



Intersection = 1



Union = 5



$$\text{Similarity} = 1/5$$


The **Hamming** similarity does not correlate with sequence content

The **Jaccard** similarity correlates with sequence content

Why?


Building an SBT

Top Down Streaming: Given an existing tree (which may be empty) and a new bloom filter, how can we build a tree?




Building an SBT

Top Down Streaming: Given an existing tree (which may be empty) and a new bloom filter, how can we build a tree?




Building an SBT

Top Down Streaming: Given an existing tree (which may be empty) and a new bloom filter, how can we build a tree?



Building an SBT

Top Down Streaming: Given an existing tree (which may be empty) and a new bloom filter, how can we build a tree?




Global Pairwise Construction: Given all bloom filters, how can we build a tree?




Building an SBT

Top Down Streaming: Given an existing tree (which may be empty) and a new bloom filter, how can we build a tree?




Global Pairwise Construction: Given all bloom filters, how can we build a tree?




Building an SBT

Top Down Streaming: Given an existing tree (which may be empty) and a new bloom filter, how can we build a tree?




Global Pairwise Construction: Given all bloom filters, how can we build a tree?




Building an SBT

Top Down Streaming: Given an existing tree (which may be empty) and a new bloom filter, how can we build a tree?




Global Pairwise Construction: Given all bloom filters, how can we build a tree?




Building an SBT

Top Down Streaming: Given an existing tree (which may be empty) and a new bloom filter, how can we build a tree?



Global Pairwise Construction: Given all bloom filters, how can we build a tree?




What are the benefits?
What are the weaknesses?

Querying a Sequence Bloom Tree

Good approximate sequence match == lots of shared kmers:


query { ATGGTTAGAATTAAACCTGGTTCTGCTAATAAACCTAGTGATGAT
kmers {
...
}



Querying a Sequence Bloom Tree

Good approximate sequence match == lots of shared kmers:


query { ATGGTTAGAATTAAACCTGGTTCTGCTAATAAACCTAGTGATGAT
kmers {
...
}



Querying a Sequence Bloom Tree

Good approximate sequence match == lots of shared kmers:


query { ATGGTTAGAATTAAACCTGGTTCTGCTAATAAACCTAGTGATGAT
kmers {
...
}



Querying a Sequence Bloom Tree

Good approximate sequence match == lots of shared kmers:

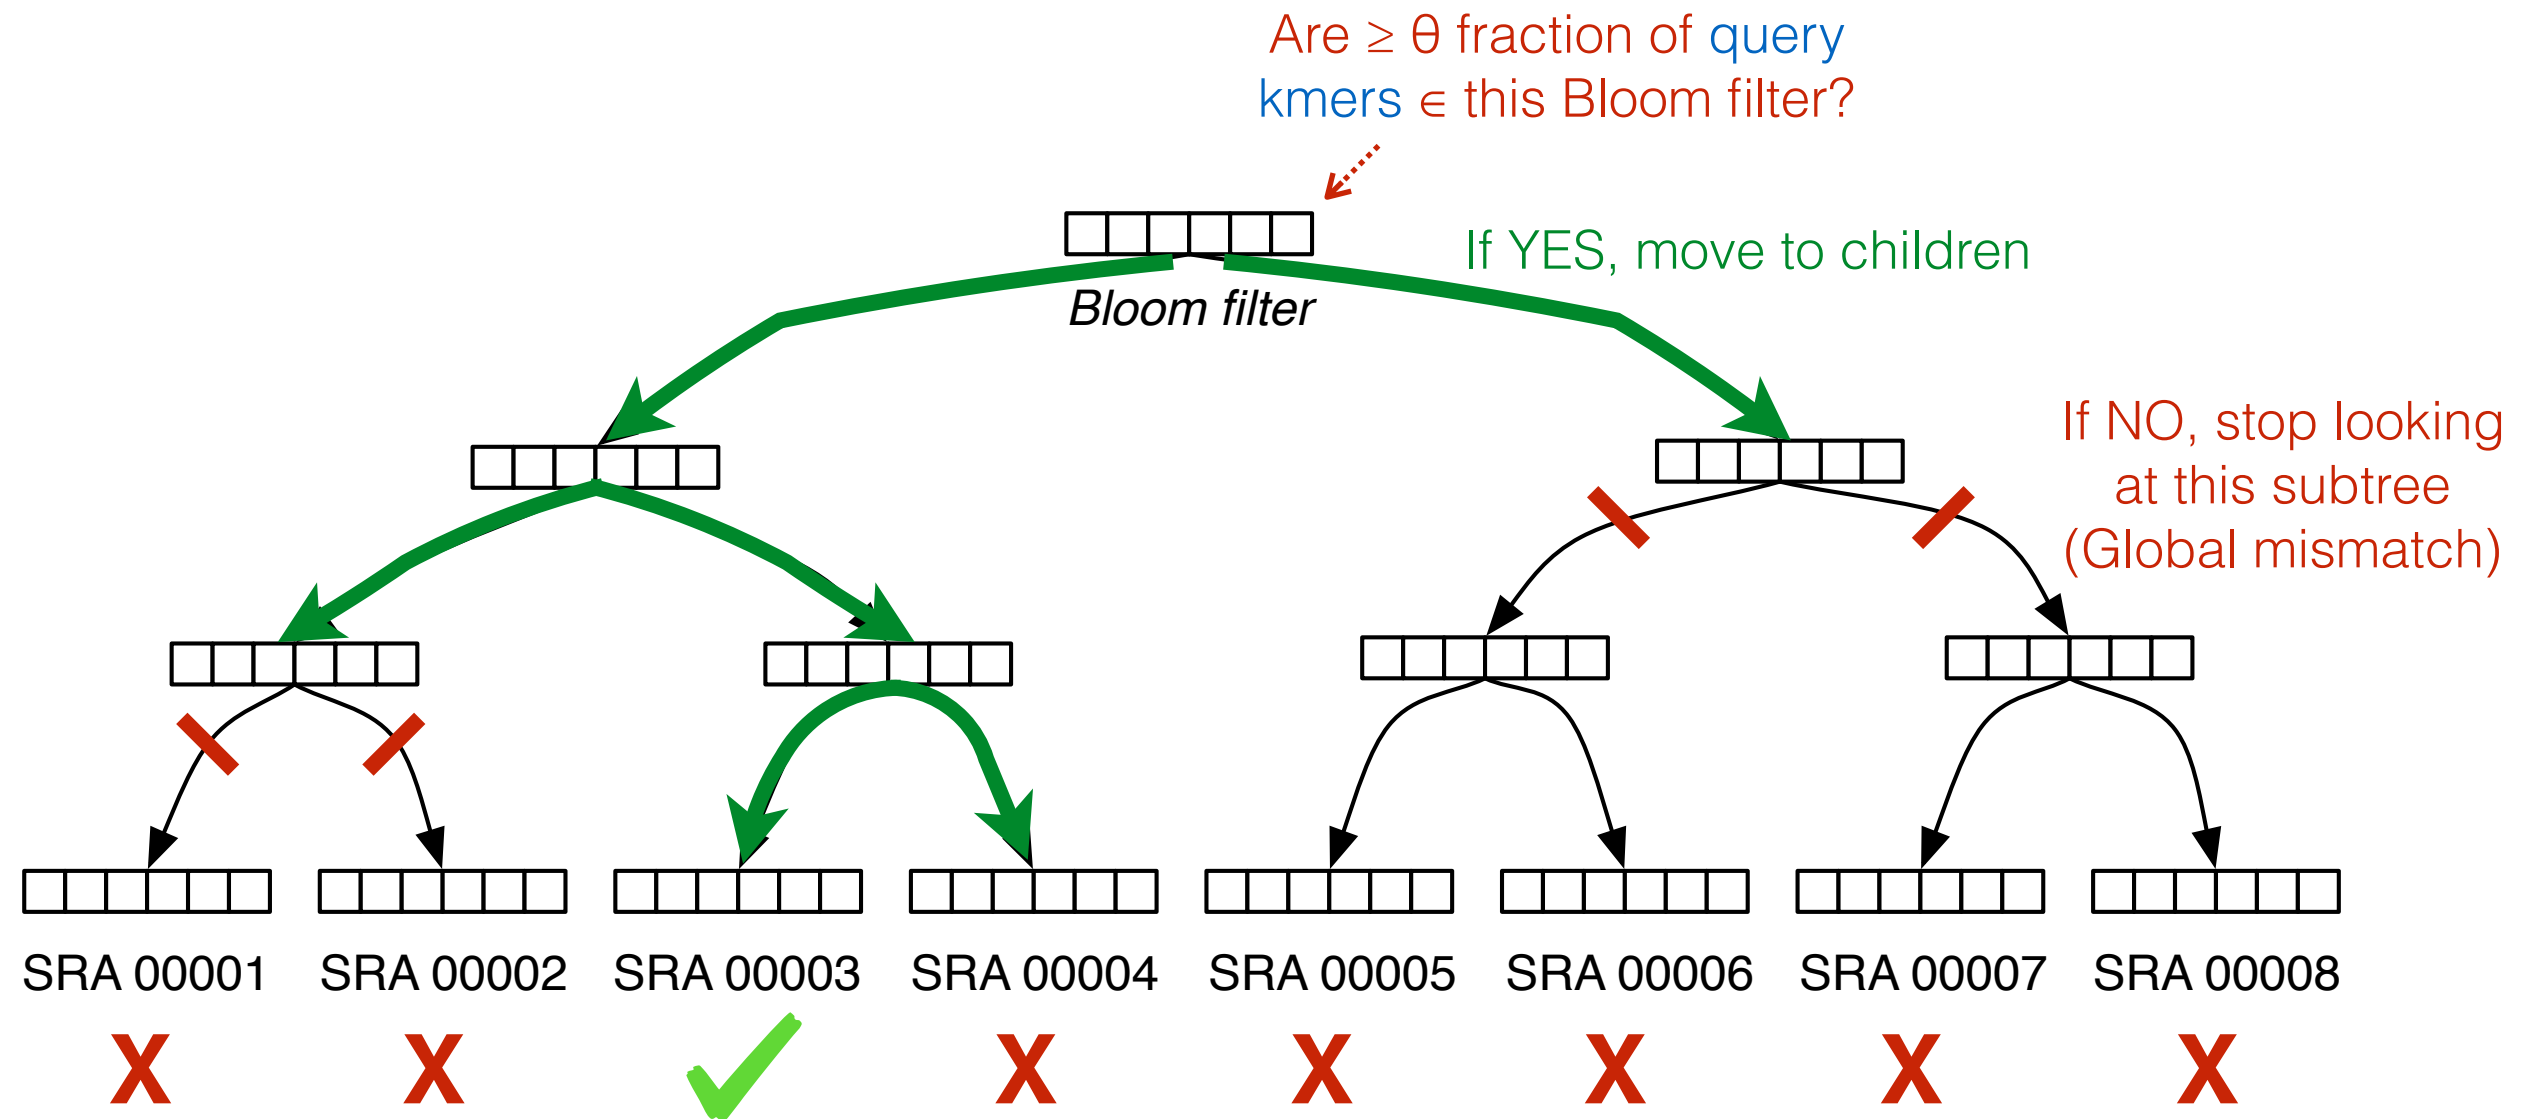
query { ATGGTTAGAATTAAACCTGGTTCTGCTAATAAACCTAGTGATGAT
kmers { ...




Querying a Sequence Bloom Tree

Good approximate sequence match == lots of shared kmers:

query { ATGGTTAGAATTAAACCTGGTTCTGCTAATAAACCTAGTGATGAT
kmers { ...




Sequence Bloom Tree is a very fast solution for containment queries




The HyperLogLog Sketch

Input items



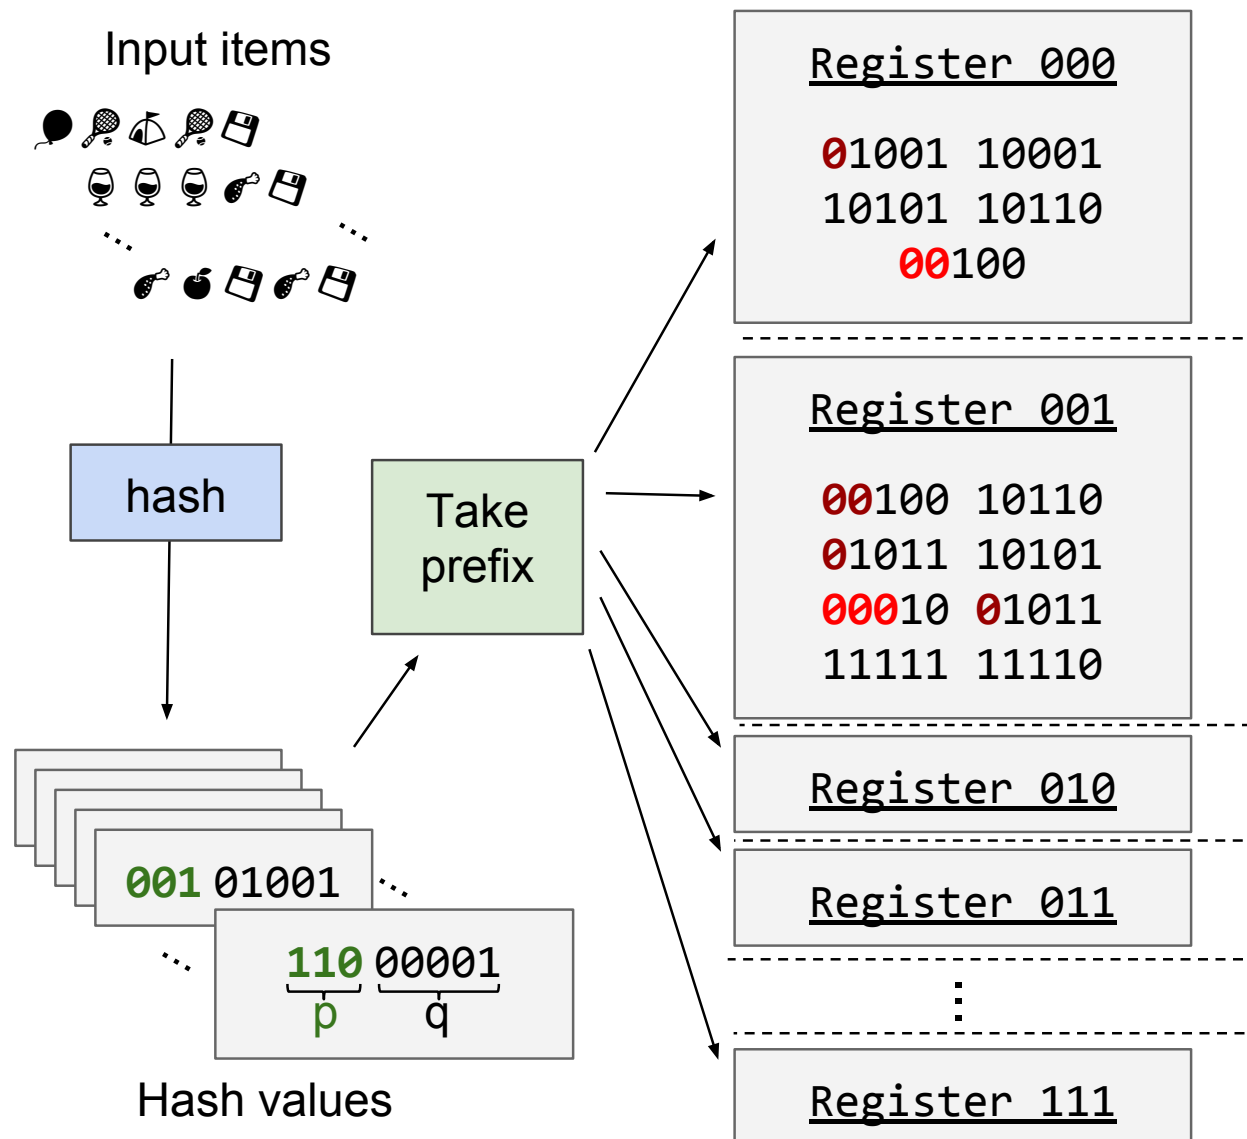
The HyperLogLog Sketch



1) Input is hashed into prefix and suffix

The HyperLogLog Sketch

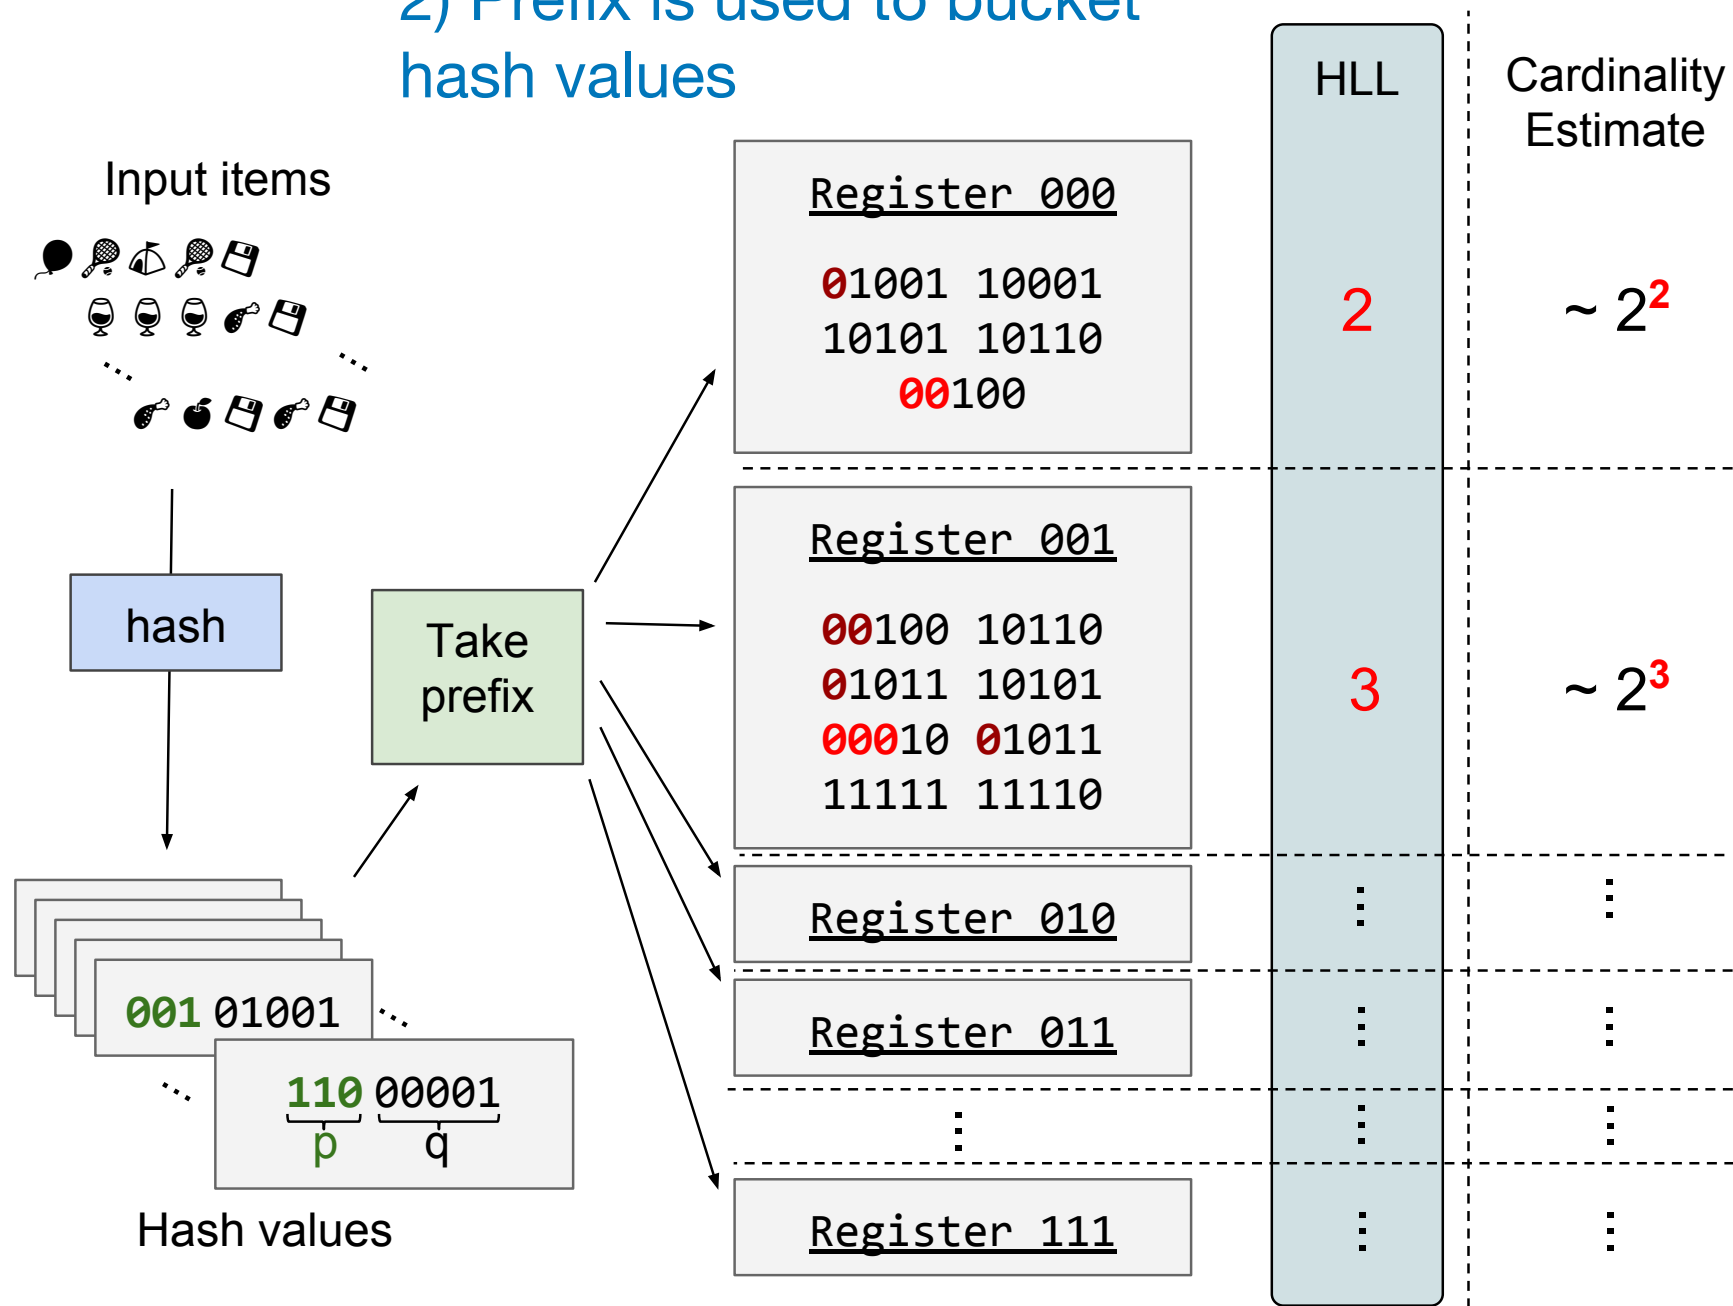
2) Prefix is used to bucket hash values



1) Input is hashed into prefix and suffix

The HyperLogLog Sketch

2) Prefix is used to bucket hash values




1) Input is hashed into prefix and suffix

3) The maximum leading zero is recorded.
Cardinality is estimated based on this value

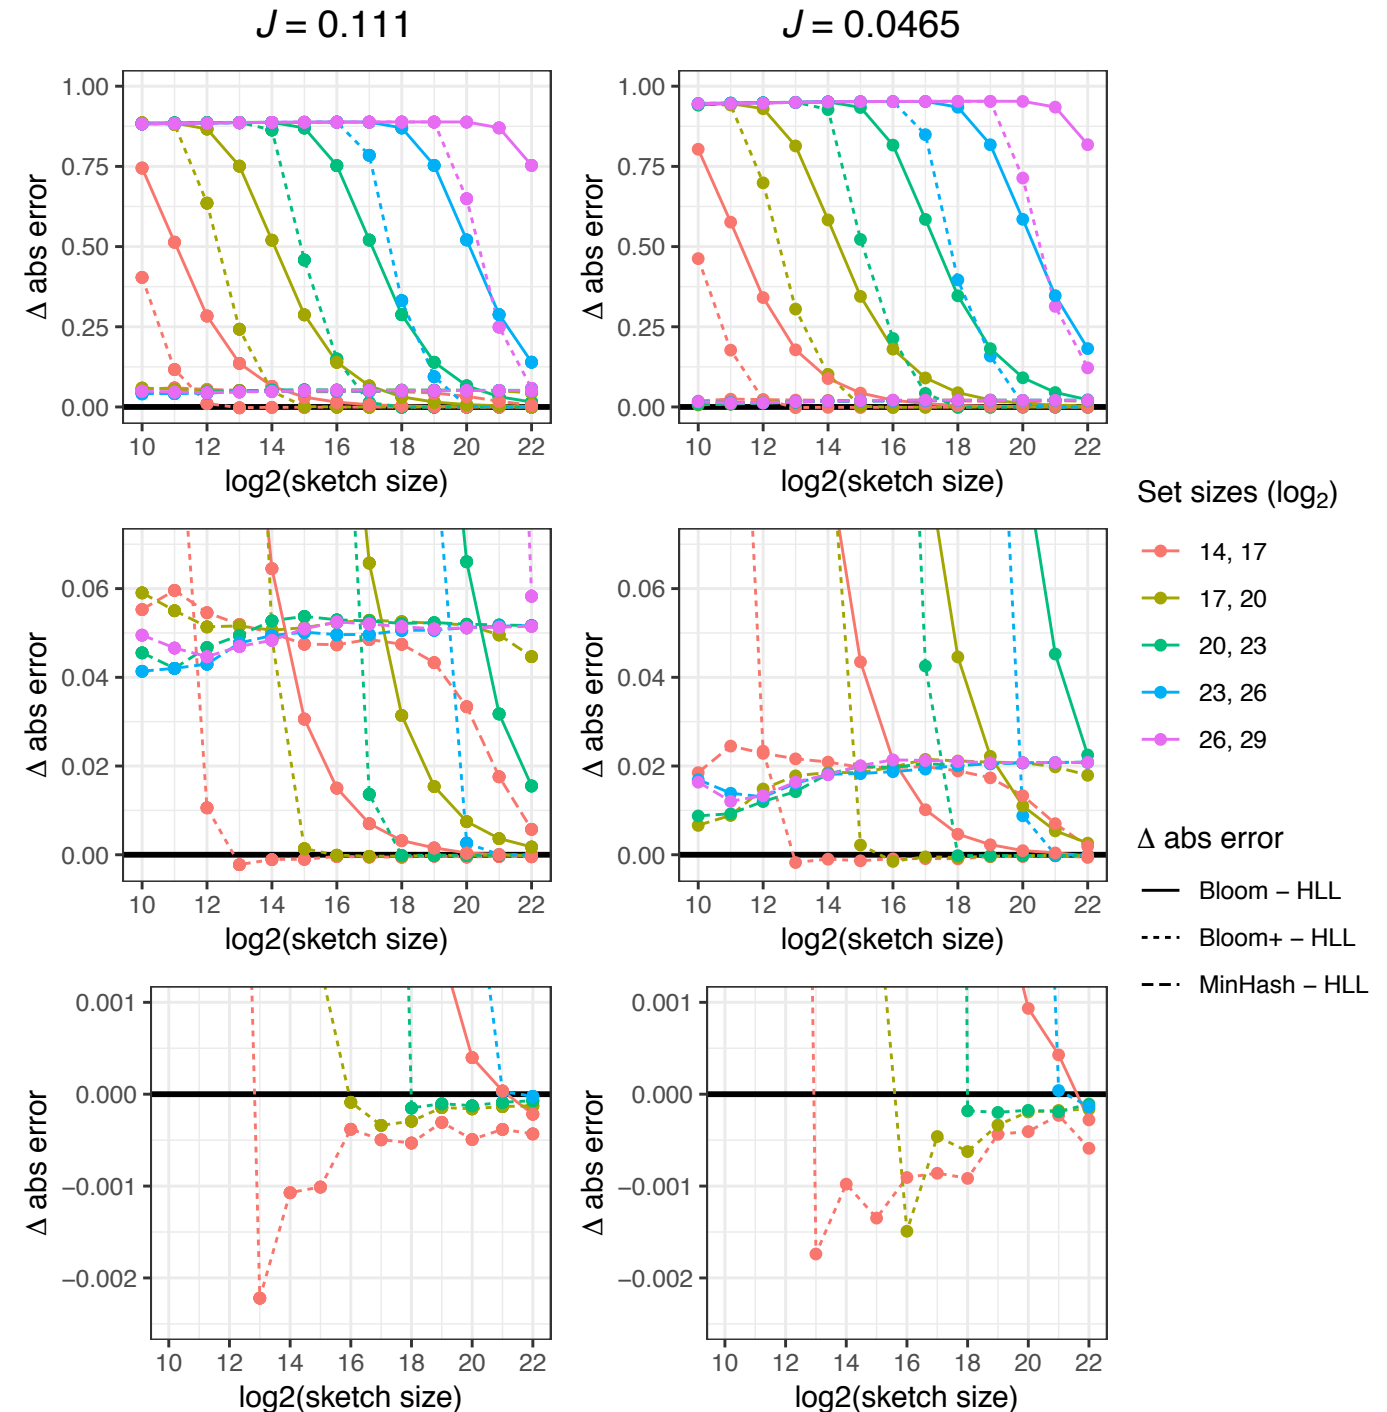
Dashing demonstrates distance difficulties for divergent data

Estimation error for computing Jaccard

Given an order of magnitude difference in data set size



HyperLogLog outperforms other sketch methods (most of the time)



The sketch size and underlying similarity can affect accuracy

The Main Take-away:

- “**Big Data**” in genomics makes conventional analysis difficult
 - Methods like Rail-RNA and recount2 try to improve efficiency through **bulk analysis**
- Sketch techniques trade accuracy for speed — and can often improve both
 - **Minhash** and **HyperLogLog** provide rapid similarity approximations
 - **Bloom Filters** provide efficient set lookup

The Main Take-away:

- “**Big Data**” in genomics makes conventional analysis difficult
 - Methods like Rail-RNA and recount2 try to improve efficiency through **bulk analysis**
- Sketch techniques trade accuracy for speed — and can often improve both
 - **Minhash** and **HyperLogLog** provide rapid similarity approximations
 - **Bloom Filters** provide efficient set lookup

Questions?

Class Projects

- Keep working on them!
- Feel free to come see me if you need help or want advice